

PHP 2550: Worksheet 7

Due: October 18th at 11:59pm

Many Analysts Recap

Summarize the data and research question posed to the teams in the paper “Many Analysts, One Data Set”. What do you notice about the difference in methodological approaches teams had? How different are the resulting estimated odds ratios? Overall, what do you think about this experiment? (~2 paragraphs)

In the paper “Many Analysts, One Data Set”, researchers applied a crowdsourcing data analysis approach to investigate the same research question: whether soccer referees are more likely to give red cards to dark-skin-toned players than to light-skin-toned players. The dataset covers 2,053 players from the top male leagues in England, Germany, France, and Spain during the 2012-2013 season, along with 3,147 referees they encountered, creating 146,028 player-referee dyads. Key variables included the players’ position, weight, height, and skin tone, and referee data such as country of origin. Additionally, each player-referee dyad provided information on how many games they interacted in, as well as the frequency of yellow and red cards issued to each player. To be noticed, the skin tone ratings, originally coded on a scale from 1 (very light) to 5 (very dark), were later standardized to a 0 to 1 scale. Additional control variables like age, club, and league which would change over players’ career are received at the time of data collection but not at the times the red cards are awarded. The same dataset was assigned to 29 independent research teams working from 13 different countries and came from various disciplinary backgrounds to investigate the same research questions. Teams are free to decide which variables they would like to include, their statistical methods to use, and their methods to handle nonindependence of players and referees. Among the 29 teams, 21 unique combinations of covariates were used, with some groups including a variety of control variables while others only used a few key variables. Additionally, statistical methods vary as well. Six teams chose to use poisson models, fifteen chose logistic models, six used linear models, and the remaining two teams used specialized methods which are classified as “miscellaneous.” Also, groups chose various methods to address the nonindependence between players and referees, including introducing fixed effects, variance components, and adding clustered standard errors.

The resulting estimated odds ratio varies among teams significantly, ranging from 0.89 to 2.93, and the median estimate is 1.31. Among the 29 teams, 20 found significant positive relationships between players' skin tone and red card issuance, other groups resulted in insignificance with no group resulting in significant negative relationship. Logistic and Poisson models generally produced higher odds ratios, with most teams finding significant effects (median ORs around 1.34–1.36), while linear models showed lower median ORs (1.21) and fewer significant results. Methods to handle nonindependence, such as fixed effects or clustered standard errors, also influenced outcomes, with teams using these methods reporting median ORs from 1.28 to 1.39. These significance variations imply how differences in variable selection and statistical model or approach would lead to various conclusions. This study introduces a really innovative approach for us to show how research is influenced by researchers' decisions and subjective choices. Even with identical dataset, different researchers could report significant various conclusions due to different choices of variables and statistical methods or analytic approaches. The paper mentions that “in some cases, authors use a particular analytic strategy because it is the one they know how to use, rather than because they have a specific rationale for using it.” This sentence does inspire us about a human side of research that we have never realized before. It implies the hidden subjective choices and personal limitations that can shape the results. By employing a crowdsourced data analysis approach, this study incorporates diverse analytic strategies and perspectives from researchers worldwide, each with unique disciplinary backgrounds. This diversity encourages transparency, highlighting the flexibility in data interpretation and strengthening the robustness and reliability of their results. The collective approach in this study implies benefits of open collaboration, as it enriches our understanding of data by drawing from a wide range of analytic viewpoints.

Answering Scientific Questions with Regression

Answer the following questions about the difference-in-differences paper you were assigned. (~1 paragraph per question)

1. What was the motivating research question? How was this translated to a scientific question and analytic approach?
2. What is the underlying model(s) used? Be as specific as possible and explain how you determined the model.

This paper uses a Difference-in-Difference approach combines with generalized linear mixed-effects model. This study has two time periods: pre-tax period (October to December 2016), and post-tax period (6 months, 12 months, and 24 months after the tax implementation). Data are collected at Philadelphia (intervention city where implements tax) and Baltimore (control city where doesn't implement tax) before and after the tax implementation. Our interested outcomes are change in beverage prices (cents per fluid ounce), change in fluid ounces of taxed and non-taxed beverages purchased per

customer, and change in total calories purchased from sweetened beverages and high-sugar foods. In addition to the main variables of interest (time period, city location, and beverage tax status), the study also includes other covariates, like income level, customer demographics (gender, race, education level, age, and frequency of store visit), and purchasing behaviors (purchased goods and total spending), to control for potential confounding. The outcomes are analyzed using a DiD approach, which compares the difference in changes over time between the two cities, aiming to isolate the effect of the tax from other factors that might have influenced beverage prices and consumption in both cities over time. Additionally, the model used generalized linear mixed-effects model which introduces random intercepts for stores to adjust for unobserved heterogeneity among stores at baseline.

explain how you determined?

3. How were the results used to answer the question and what was the conclusion?

- The results in the “Change in Beverage Price” can be used to answer the question of In the paper, the research mentions that “there was a 2.06 cents per fl oz (95% CI, 1.75 to 2.38 cents per fl oz; $p < .001$) increase for taxed beverages in Philadelphia compared with Baltimore, an increase of 33%, indicating a 137.3% pass-through of tax.” This implies that stores not only passed the 1.5-cent-per-ounce tax onto customers, but also charge even more than the required tax. Since Baltimore, the control city, does not show price increases and non-taxed beverages in Philadelphia also does not exhibit this price increase, the price increases in Philadelphia would be the result of the tax implementation. Also, the price increases were consistent across income levels, with no significant difference between low- and high-income neighborhoods. By using DiD approach combined with generalized linear mixed-effects models, researchers are able to compare the price change on taxed and non-taxed beverages in Philadelphia and Baltimore before and after the tax was implemented. Also, the adding random intercepts across stores is able to adjust for differences in store-specific factors. From the results, we can conclude that the tax implement significantly increase the price of taxed beverage overtime, and the tax are overly passed on to the customers which indicates the direct and substantial impact of tax on beverage prices.
- The results in the “Change in Volume of Beverage Purchased” can be used to answer the question of The paper mentions that “there was a 6.12-fl oz decline (95% CI, -9.88 to -2.37 fl oz; $P < .001$), or a 41.9% decrease, in the ounces of taxed beverages purchased per person in Philadelphia compared with Baltimore.” This refers to a significant drop in the purchase of taxed beverages. This significant drop was driven mainly by a reduction in the purchase of sugar-sweetened beverages (SSBs), with a 6.17 fluid ounce per person reduction, corresponding to a 47.3% decrease. Also, the paper said there was no significant change on the volume of non-taxed beverages in Philadelphia, indicates the implementation of tax would only increase

the sweet beverages. In addition, the analysis showed that the tax's effects were more pronounced in low-income neighborhoods and lower education levels, with purchases dropped by around 40%. From their results, we can conclude that the tax implementation in Philadelphia significantly reduced the consumption amount of taxed beverages where the effect is more pronounced on vulnerable population. From this, we can see that tax implementation is able to help reducing sugar consumptions and addressing the health disparity problem caused by excessive sugar consumption.

- The results in the “Changes in Calories and Spending on Beverages and High-Sugar Foods” can be used to answer the question of The paper mentions that “there was a 69-calorie decrease (95% CI, −132 to −5 calories; $P = .04$) in the total calories purchased from SSBs and high-sugar foods combined, a 22.6% decline.” In addition, “the grams of sugar from these items declined by 19.9 g (95% CI, −31.7 to −8.2 g; $P = .002$), or 34.1% per person.” These results indicate that the tax implementation significantly reduce calorie and sugar intake. There wasn't change on people's spending with their shopping post tax, but the frequency of neighboring counties purchasing increases slightly that people chose to make purchases in other counties to avoid tax. This reduction was more noticeable among low-income neighborhoods and people with lower levels of education as well. From these, we can conclude tax significantly decrease people's calorie and sugar intake from sweetened goods, especially strong for people with lower-income and lower levels of education. However, this decrease does not influence their total spendings.
4. Overall, how do you evaluate this paper? Think about strengths and weaknesses of the approach and any remaining questions you have.

Model Evaluation Example

These questions are on the paper ‘Predicting lung cancer prior to surgical resection in patients with lung nodules’ by Deppen et al. This paper introduces a model called TREAT that is currently used in practice to predict lung cancer.

1. Compare the Mayo model to the TREAT model in terms of the initial goals of building the model, the population the training data represented, the variables included, and the resulting model. (~2 paragraphs)
2. What measures or visuals were used to evaluate the models? How do we interpret these? Why do you think these measures were chosen for comparison? (1 paragraph)
3. What were some limitations that the paper addressed? (1 paragraph)

Model Building Practice

Read the [NEJM editorial](#) to understand the background of developing the equations that are used to calculate the estimated glomerular filtration rate (eGFR). Then, load in the data `baseseg.csv` and run the pre-processing below. The data contains the following variables.

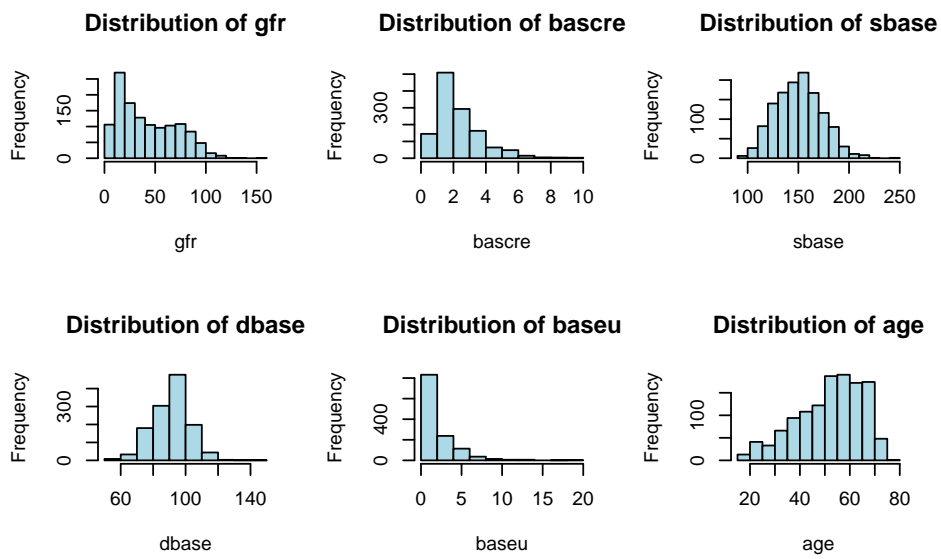

1. Base serum Creatinine (`bascre`)
2. Systolic blood pressure (`sbase`)
3. Diastolic blood pressure (`dbase`)
4. Urine protein (`baseu`)
5. Age (`age`)
6. Sex (Sex = 1 if male; = 0 if female)
7. Indicator if African-American (`black`)
8. Measured glomerular filtration rate (`gfr`)

In this worksheet, we will build a model to calculate the eGFR and practice our model evaluation skills.

1. Conduct a brief exploratory data analysis (EDA). Check the distributions of the variables and consider whether transformations are necessary. Hints: Log transformations or polynomial transformations may be helpful.
2. Build a linear regression model with appropriate variable selection. Check the model assumptions using model diagnostics. You may also consider including interaction terms.
3. Evaluate the performance of the model using evaluation measures. Using your evaluation and your estimated model, comment on how useful you expect the model to be in practice.
4. Evaluate the performance of the model now between race populations. In particular, compare the measured and estimated GFR using mean squared error (MSE), bias, and the percentage of estimates within 10% and 30% of the measured GFR (P_{10} (%) and P_{30} (%), respectively). Visualize the comparison of the measured and estimated GFR.
5. Repeat steps 1-4 but remove the race variable (`black`) from consideration. Interpret your results and relate them back to the discussion in the editorial.
6. Last, write a non-technical summary of one of your models and its evaluation (1 paragraph) for a clinical audience.

1

Figure 1: Box Plot of gfr by Sex



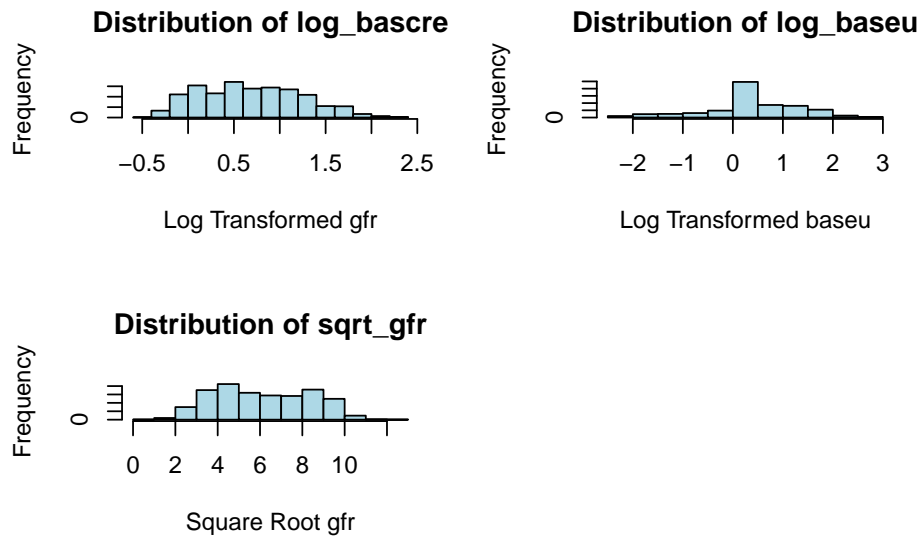


Table 1: Summary Table by Race

Characteristic	Original Variables		Transformed Variables	
	Non-Black, N = 1,135	Black, N = 114	Non-Black, N = 1,135	Black, N = 114
age	52 (13)	52 (11)	52 (13)	52 (11)
sex				
Female	425 (37%)	41 (36%)	425 (37%)	41 (36%)
Male	710 (63%)	73 (64%)	710 (63%)	73 (64%)
Measured glomerular filtration rate	43 (29)	38 (21)	6.16 (2.27)	5.90 (1.77)
Base serum Creatinine	2.31 (1.41)	2.55 (1.32)	0.67 (0.56)	0.83 (0.43)
Systolic blood pressure	154 (23)	139 (24)	154 (23)	139 (24)
Diastolic blood pressure	93 (11)	89 (12)	93 (11)	89 (12)
Urine protein	1.91 (2.32)	0.94 (1.99)	0.41 (0.89)	-0.42 (1.07)

¹ Mean (SD) for continuous; n (%) for categorical

Figure 2: Correlation Matrix


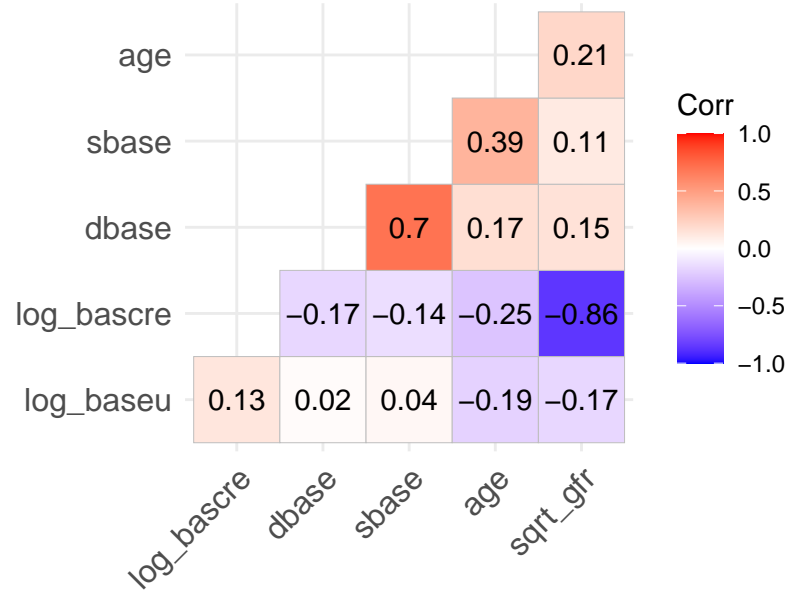



Table 2: Coefficient Estimate of the Best Model

	Estimate	Standard Error	T Statistics	P Value
(Intercept)	6.1055730	0.5007233	12.1935066	0.0000000

Table 2: Coefficient Estimate of the Best Model (*continued*)

	Estimate	Standard Error	T Statistics	P Value
dbase	0.0191699	0.0049231	3.8938412	0.0001040
age	0.0039419	0.0040693	0.9686918	0.3328888
sex1	1.7481928	0.3045048	5.7411011	0.0000000
black1	-0.7740926	0.2778949	-2.7855586	0.0054253
log_bascre	-1.2169126	0.4928199	-2.4692844	0.0136729
log_baseu	-1.0255112	0.3012030	-3.4047172	0.0006836
dbase:log_bascre	-0.0227782	0.0051811	-4.3963883	0.0000120
dbase:log_baseu	0.0044291	0.0028808	1.5374605	0.1244369
age:sex1	-0.0166141	0.0050867	-3.2661645	0.0011201
age:log_baseu	0.0046374	0.0026505	1.7496454	0.0804281
sex1:black1	0.5058117	0.2239153	2.2589424	0.0240609
sex1:log_bascre	-0.4557435	0.1185290	-3.8449967	0.0001267
black1:log_bascre	0.7106491	0.2451180	2.8992126	0.0038072
log_bascre:log_baseu	0.2703598	0.0729011	3.7085843	0.0002177



3

The MSE is: 1.167612


Table 3: Performance Metrics by Race

black	MSE	Bias	P10	P30
0	1.158960	0.0370461	49.57265	84.18803
1	1.312231	-0.3454632	50.00000	92.85714

The MAE is: 0.7988623

[1] 160.8011

4



5




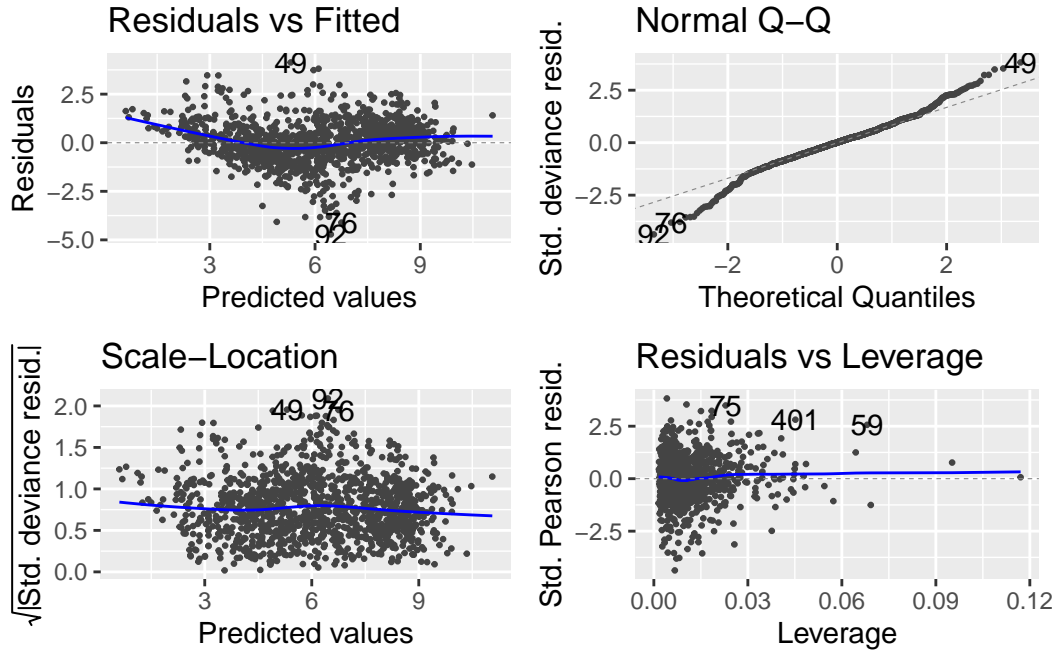
Table 4: Summary Table by Race

Characteristic	Original Variables		Transformed Variables		Race Removed
	Non-Black, N = 1,135	Black, N = 114	Non-Black, N = 1,135	Black, N = 114	N = 1,249
age	52 (13)	52 (11)	52 (13)	52 (11)	52 (13)
sex					
Female	425 (37%)	41 (36%)	425 (37%)	41 (36%)	466 (37%)
Male	710 (63%)	73 (64%)	710 (63%)	73 (64%)	783 (63%)
Measured glomerular filtration rate	43 (29)	38 (21)	6.16 (2.27)	5.90 (1.77)	6.13 (2.23)
Base serum Creatinine	2.31 (1.41)	2.55 (1.32)	0.67 (0.56)	0.83 (0.43)	0.69 (0.55)
Systolic blood pressure	154 (23)	139 (24)	154 (23)	139 (24)	152 (23)
Diastolic blood pressure	93 (11)	89 (12)	93 (11)	89 (12)	93 (12)
Urine protein	1.91 (2.32)	0.94 (1.99)	0.41 (0.89)	-0.42 (1.07)	0.33 (0.94)

¹ Mean (SD) for continuous; n (%) for categorical

Table 5: Coefficient Estimate of the Best Model without Race

	Estimate	Standard Error	T Statistics	P Value
(Intercept)	5.5337283	0.5789457	9.5582862	0.0000000
dbase	0.0256797	0.0059645	4.3054386	0.0000180
age	0.0025477	0.0041172	0.6188018	0.5361610
sex1	2.5098270	0.5811383	4.3188121	0.0000169
log_bascre	-1.0472561	0.4906968	-2.1342224	0.0330209
log_baseu	-1.0739182	0.3022504	-3.5530743	0.0003951
dbase:sex1	-0.0089447	0.0056697	-1.5776354	0.1149053
dbase:log_bascre	-0.0241935	0.0051649	-4.6842460	0.0000031
dbase:log_baseu	0.0051956	0.0029065	1.7875634	0.0740914
age:sex1	-0.0147942	0.0051655	-2.8640189	0.0042537
age:log_baseu	0.0042962	0.0026609	1.6145709	0.1066589
sex1:log_bascre	-0.4326624	0.1192452	-3.6283428	0.0002969
log_bascre:log_baseu	0.2546948	0.0723081	3.5223536	0.0004433




The MSE is: 1.167296

The MAE is: 0.8021315

Table 6: Performance Metrics by Race

black	MSE	Bias	P10	P30
0	1.153581	0.0496263	48.71795	84.18803
1	1.396532	-0.5743704	57.14286	85.71429



Appendix

```
suppressPackageStartupMessages(library(tidyverse))
suppressPackageStartupMessages(library(moments))
suppressPackageStartupMessages(library(stats))
suppressPackageStartupMessages(library(kableExtra))
suppressPackageStartupMessages(library(dplyr))
suppressPackageStartupMessages(library(gtsummary))
suppressPackageStartupMessages(library(gt))
suppressPackageStartupMessages(library(corrplot))
suppressPackageStartupMessages(library(ggcorrplot))
suppressPackageStartupMessages(library(MASS))
suppressPackageStartupMessages(library(ggplot2))
suppressPackageStartupMessages(library(ggfortify))
suppressPackageStartupMessages(library(caret))
```

```

suppressPackageStartupMessages(library(kableExtra))

kidney_df <- read.csv("baseseg.csv")
kidney_df <- kidney_df %>%
  dplyr::select(gfr, bascre, sbase, dbase, baseu, AGE, SEX, black) %>%
  rename(sex = SEX, age = AGE) %>%
  na.omit()

kidney_df$black <- as.factor(kidney_df$black)
kidney_df$sex <- as.factor(kidney_df$sex)
ggplot(kidney_df) +
  geom_boxplot(aes(x = sex, y = gfr, fill = sex)) +
  ggtitle("Figure 1: Box Plot of gfr by Sex") +
  theme_minimal()

par(mfrow = c(2, 3))
hist(kidney_df$gfr, main = "Distribution of gfr", xlab = "gfr", col = "lightblue")
hist(kidney_df$bascre, main = "Distribution of bascre", xlab = "bascre", col = "lightblue")
hist(kidney_df$sbase, main = "Distribution of sbase", xlab = "sbase", col = "lightblue")
hist(kidney_df$dbase, main = "Distribution of dbase", xlab = "dbase", col = "lightblue")
hist(kidney_df$baseu, main = "Distribution of baseu", xlab = "baseu", col = "lightblue")
hist(kidney_df$age, main = "Distribution of age", xlab = "age", col = "lightblue")

# # Calculate skewness
# skewness_value <- skewness(kidney_df[, 1:6])
# knitr::kable(skewness_value,
#               caption = "Variable Skewness Value",
#               col.names = c("Variable", "Skewness")) %>%
#   kable_styling(latex_options = "HOLD_position")

kidney_df$log_bascre <- log(kidney_df$bascre)
kidney_df$log_baseu <- ifelse(kidney_df$baseu <= 0.1, kidney_df$baseu, log(kidney_df$baseu))
kidney_df$sqrt_gfr <- sqrt(kidney_df$gfr)
kidney_df$log_gfr <- log(kidney_df$gfr)

par(mfrow = c(2, 2))
hist(kidney_df$log_bascre, main = "Distribution of log_bascre", xlab = "Log Transformed gfr", col = "lightblue")
hist(kidney_df$log_baseu, main = "Distribution of log_baseu", xlab = "Log Transformed baseu", col = "lightblue")
hist(kidney_df$sqrt_gfr, main = "Distribution of sqrt_gfr", xlab = "Square Root gfr", col = "lightblue")
# hist(kidney_df$log_gfr, main = "Distribution of log_gfr", xlab = "Log Transformed gfr", col = "lightblue")
# Creating first summary table (original)

```

```

kidney_tbl <- kidney_df %>%
  dplyr::select(black, age, sex, gfr, bascre, sbase, dbase, baseu) %>%
  mutate(black = recode(black, `0` = "Non-Black", `1` = "Black"),
         sex = recode(sex, `0` = "Female", `1` = "Male")) %>%
  tbl_summary(by = black,
             label = list(gfr ~ "Measured glomerular filtration rate",
                          bascre ~ "Base serum Creatinine",
                          sbase ~ "Systolic blood pressure",
                          dbase ~ " Diastolic blood pressure",
                          baseu ~ "Urine protein"
                          ),
             statistic = all_continuous() ~ "{mean} ({sd})") %>%
  modify_spanning_header(update = all_stat_cols() ~ "**Black**") %>%
  modify_footnote(update = all_stat_cols() ~ "Mean (SD) for continuous; n (%) for categori
  bold_labels()

# Creating second summary table (transformed)
kidney_tbl_transformed <- kidney_df %>%
  dplyr::select(black, age, sex, sqrt_gfr, log_bascre, sbase, dbase, log_baseu) %>%
  rename(gfr=sqrt_gfr, bascre=log_bascre, baseu=log_baseu) %>%
  mutate(black = recode(black, `0` = "Non-Black", `1` = "Black"),
         sex = recode(sex, `0` = "Female", `1` = "Male")) %>%
  tbl_summary(by = black,
             label = list(gfr ~ "Measured glomerular filtration rate",
                          bascre ~ "Base serum Creatinine",
                          sbase ~ "Systolic blood pressure",
                          dbase ~ " Diastolic blood pressure",
                          baseu ~ "Urine protein"
                          ),
             statistic = all_continuous() ~ "{mean} ({sd})") %>%
  modify_spanning_header(update = all_stat_cols() ~ "**Black**") %>%
  modify_footnote(update = all_stat_cols() ~ "Mean (SD) for continuous; n (%) for categori
  bold_labels()

tbl_combined1 <- tbl_merge(
  tbls = list(kidney_tbl, kidney_tbl_transformed),
  tab_spanner = c("**Original Variables**", "**Transformed Variables**")
)

tbl_combined1 <- tbl_combined1 %>%
  as_kable_extra(booktabs = TRUE, caption = "Summary Table by Race",

```

```

        longtable = TRUE, linesep = "") %>%
kableExtra::kable_styling(font_size = 10,
        latex_options = c("repeat_header", "HOLD_position", "scale_dow
column_spec(1, width = "3cm") %>%
column_spec(2, width = "2cm") %>%
column_spec(3, width = "2cm") %>%
column_spec(4, width = "2cm") %>%
column_spec(5, width = "2cm") %>%
row_spec(0, bold = TRUE)

tbl_combined1
cor_m <- cor(kidney_df[, -c(1, 2, 5, 7, 8, 12)])
variable_order <- c("log_baseu", "log_bascre", "dbase", "sbase", "age", "sqrt_gfr")
r_reordered <- cor_m[variable_order, variable_order]

ggcorrplot(r_reordered,
            hc.order = TRUE,
            type = "lower",
            lab = TRUE) +
  ggtitle("Figure 2: Correlation Matrix") +
  theme(plot.title = element_text(hjust = 0.5, size=14),
        axis.text.x = element_text(size = 12),
        axis.text.y = element_text(size = 12))
full_model <- glm(sqrt_gfr ~ dbase + age + sex + black + log_bascre + log_baseu +
  dbase:age + dbase:sex + dbase:black + dbase:log_bascre + dbase:log_baseu +
  age:sex + age:black + age:log_bascre + age:log_baseu +
  sex:black + sex:log_bascre + sex:log_baseu +
  black:log_bascre + black:log_baseu + log_bascre:log_baseu,
  data = kidney_df)

step_model <- step(full_model, direction = "backward", trace = 0)

autoplot(step_model, size=0.5)

summary_model <- as.data.frame(summary(step_model)$coefficient)
colnames(summary_model) <- c("Estimate", "Standard Error", "T Statistics", "P Value")

summary_model %>%
  kbl(booktabs = TRUE, caption = "Coefficient Estimate of the Best Model",
      longtable = TRUE, linesep = "") %>%
  kable_styling(font_size = 10,

```



```

    latex_options = c("repeat_header", "HOLD_position", "scale_down"))

autoplot(step_model, size=0.5)
set.seed(2550)
index <- createDataPartition(kidney_df$gfr, p = 0.8, list = FALSE)
train <- kidney_df[index, ]
test <- kidney_df[-index, ]
predictions <- predict(step_model, newdata = test)
mse <- mean((predictions - test$sqrt_gfr)^2)
mae <- mean(abs(predictions - test$sqrt_gfr))

cat("The MSE is:", mse, "\n")
cat("The MAE is:", mae, "\n")

# Back-transform the predictions
test$predictions_transformed <- predictions^2

# Compute the Mean Squared Error (MSE)
# Note: Able to directly compare gfr to predictions^2
MSE <- mean((test$gfr - test$predictions_transformed)^2)

MSE

# Split data by race and calculate performance metrics
test$estimated_gfr <- predict(step_model, newdata = test)
performance_by_race <- test %>%
  group_by(black) %>%
  summarise(MSE = mean((estimated_gfr - sqrt_gfr)^2),
            Bias = mean(estimated_gfr - sqrt_gfr),
            P10 = mean(abs(estimated_gfr - sqrt_gfr) / sqrt_gfr <= 0.10) * 100,
            P30 = mean(abs(estimated_gfr - sqrt_gfr) / sqrt_gfr <= 0.30) * 100)
performance_by_race %>%
  kbl(caption = "Performance Metrics by Race") %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed", "responsive"))

# Create a scatter plot to compare measured and estimated GFR
ggplot(test, aes(x = sqrt_gfr, y = estimated_gfr, color = black)) +
  geom_point(alpha=0.8) +
  geom_abline(intercept = 0, slope = 1, linetype = "solid", color = "black", size=1) + #
  geom_smooth(method = "lm", linetype = "dashed", se = FALSE, alpha=0.8) +

```

```

scale_color_manual(values = c("0" = "red", "1" = "blue"), # Assign colors manually
                    labels = c("0" = "Non-black", "1" = "Black"), # Rename legend labels
                    name = "Race") +
labs(title = "Comparison of Measured and Estimated GFR by Race",
     x = "Measured GFR",
     y = "Estimated GFR") +
theme_minimal()
cols <- c("#F76D5E", "#FFFFBF")
#, "#72D8FF"
# Basic density plot in ggplot2
cols <- c("Black" = "#F76D5E", "Non-Black" = "#FFFFBF", "All" = "#72D8FF")
ggplot() +
  geom_density(data = kidney_df, aes(x = sqrt_gfr, fill = "All"),
              alpha = 1, color = "#72D8FF") + # Entire dataset
  geom_density(data = kidney_df %>% filter(black == 1),
              aes(x = sqrt_gfr, fill = "Black"),
              alpha = 0.5, color = "#F76D5E") + # Black group
  geom_density(data = kidney_df %>% filter(black == 0),
              aes(x = sqrt_gfr, fill = "Non-Black"),
              alpha = 0.5, color = "#FFFFBF") + # Non-Black group
scale_fill_manual(values = cols) +
labs(title = "Density Plot of sqrt_gfr by Race",
     x = "sqrt_gfr",
     y = "Density",
     fill = "Group") +
theme_minimal()
# Creating third summary table (transformed)
kidney_tbl_nb <- kidney_df %>%
  dplyr::select(sex, age, sqrt_gfr, log_bascre, sbase, dbase, log_baseu) %>%
  rename(gfr=sqrt_gfr, bascre=log_bascre, baseu=log_baseu) %>%
  mutate(sex = recode(sex, `0` = "Female", `1` = "Male")) %>%
  tbl_summary(label = list(gfr ~ "Measured glomerular filtration rate",
                          bascre ~ "Base serum Creatinine",
                          sbase ~ "Systolic blood pressure",
                          dbase ~ "Diastolic blood pressure",
                          baseu ~ "Urine protein"
                          ),
             statistic = all_continuous() ~ "{mean} ({sd})") %>%
  modify_footnote(update = all_stat_cols() ~ "Mean (SD) for continuous; n (%) for categorical")
  bold_labels()

```

```

tbl_combined2 <- tbl_merge(
  tbls = list(kidney_tbl, kidney_tbl_transformed, kidney_tbl_nb),
  tab_spanner = c("**Original Variables**", "**Transformed Variables**", "**Race Removed**")
)

tbl_combined2 <- tbl_combined2 %>%
  as_kable_extra(booktabs = TRUE, caption = "Summary Table by Race",
    longtable = TRUE, linesep = "") %>%
  kableExtra::kable_styling(font_size = 10,
    latex_options = c("repeat_header", "HOLD_position", "scale_down"))
  column_spec(1, width = "3cm") %>%
  column_spec(2, width = "2cm") %>%
  column_spec(3, width = "2cm") %>%
  column_spec(4, width = "2cm") %>%
  column_spec(5, width = "2cm") %>%
  column_spec(6, width = "2cm") %>%
  row_spec(0, bold = TRUE)

tbl_combined2 <- tbl_combined2 %>%
  row_spec(5, background = "#e0ecf4") %>%
  row_spec(6, background = "#e0ecf4") %>%
  row_spec(9, background = "#e0ecf4")

tbl_combined2
model_norace <- glm(sqrt_gfr ~ sbase + dbase + age + sex + log_bascre + log_baseu +
  dbase:age + dbase:sex + dbase:log_bascre + dbase:log_baseu + age:sex +
  age:log_bascre + age:log_baseu + sex:log_bascre + sex:log_baseu +
  log_bascre:log_baseu,
  data = kidney_df)

step_model_norace <- step(model_norace, direction = "backward", trace = 0)

summary_model_norace <- as.data.frame(summary(step_model_norace)$coefficient)
colnames(summary_model_norace) <- c("Estimate", "Standard Error", "T Statistics", "P Value")

summary_model_norace %>%
  kbl(booktabs = TRUE, caption = "Coefficient Estimate of the Best Model without Race",
    longtable = TRUE, linesep = "") %>%
  kable_styling(font_size = 10,
    latex_options = c("repeat_header", "HOLD_position", "scale_down"))

```

```

autoplot(step_model_norace, size=0.5)
set.seed(2550)
index_norace <- createDataPartition(kidney_df$gfr, p = 0.8, list = FALSE)
train_norace <- kidney_df[index_norace, ]
test_norace <- kidney_df[-index_norace, ]

predictions_norace <- predict(step_model_norace, newdata = test_norace)
mse_norace <- mean((predictions_norace - test_norace$sqrt_gfr)^2)
mae_norace <- mean(abs(predictions_norace - test_norace$sqrt_gfr))

cat("The MSE is:", mse_norace, "\n")
cat("The MAE is:", mae_norace, "\n")
# Split data by race and calculate performance metrics
test_norace$estimated_gfr <- predict(step_model_norace, newdata = test_norace)
performance_by_race_norace <- test_norace %>%
  group_by(black) %>%
  summarise(MSE = mean((estimated_gfr - sqrt_gfr)^2),
            Bias = mean(estimated_gfr - sqrt_gfr),
            P10 = mean(abs(estimated_gfr - sqrt_gfr) / sqrt_gfr <= 0.10) * 100,
            P30 = mean(abs(estimated_gfr - sqrt_gfr) / sqrt_gfr <= 0.30) * 100)
performance_by_race_norace %>%
  kbl(caption = "Performance Metrics by Race") %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed", "responsive"))

# Create a scatter plot to compare measured and estimated GFR
ggplot(test_norace, aes(x = sqrt_gfr, y = estimated_gfr, color = black)) +
  geom_point(alpha=0.8) +
  geom_abline(intercept = 0, slope = 1, linetype = "solid", color = "black", size=1) + #
  geom_smooth(method = "lm", linetype = "dashed", se = FALSE, alpha=0.8) +
  scale_color_manual(values = c("0" = "red", "1" = "blue"), # Assign colors manually
                    labels = c("0" = "Non-black", "1" = "Black"), # Rename legend labels
                    name = "Race") +
  labs(title = "Comparison of Measured and Estimated GFR by Race",
       x = "Measured GFR",
       y = "Estimated GFR") +
  theme_minimal()

```